


# TurboNet: Faithfully Emulating Networks with Programmable Switches

Jiamin Cao (Tsinghua University)

Yu Zhou (Tsinghua University), Ying Liu (Tsinghua University),  
Mingwei Xu (Tsinghua University), Yongkai Zhou (UnionPay)



# To conduct network experiments



Simulators ?

Emulators ?

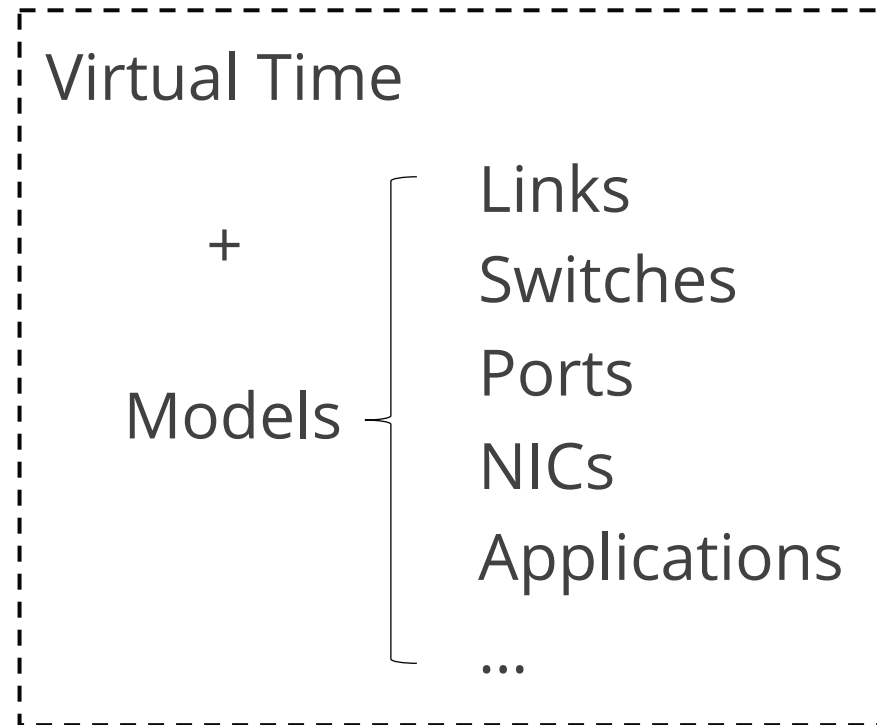
Test-beds ?

## Network experiments platforms should satisfy:

- ❑ **High Fidelity:** like realistic hardware environment
- ❑ **High Flexibility:** topologies, properties (link delay, packet losses, link bandwidth, ...)
- ❑ **High Scalability:** easily scale to meet increasing network bandwidth and sizes

# Simulators

Use CPU to model networks as simulated events



 **NS-3**  
NETWORK SIMULATOR

 **OPNET**<sup>®</sup>  
Making Networks and Applications Perform<sup>®</sup>

 **OMNeT++**  
Discrete Event Simulator

***Fidelity Concerns***

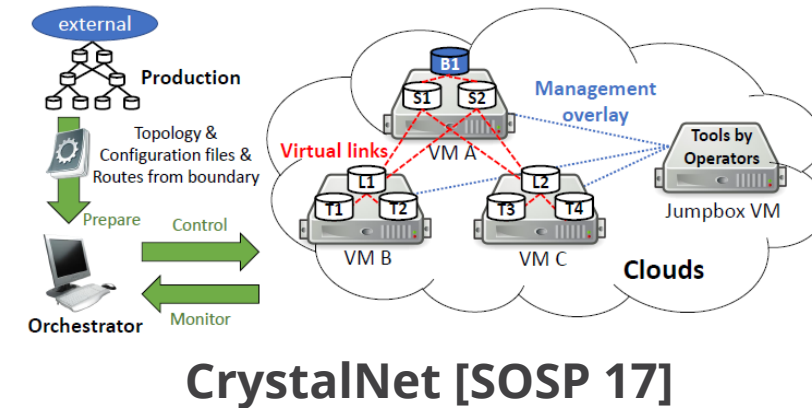
# Emulators

Run the same code on CPU as real platforms with interactive network traffic




**Cannot provide accurate performance results for large and high-speed networks**

- 10Gbps networks?
- 100Gbps networks?



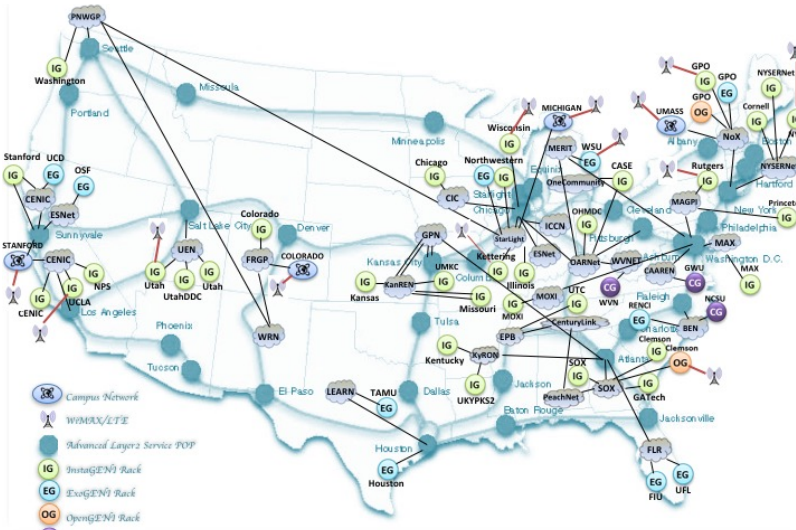
`> sudo mn`



Mininet

***Fidelity & Scalability Concerns***

# Test-Beds



GENI



CERNET2




PlanetLab

## Experiments with real-world settings

- 😊 Can guarantee fidelity
- 😞 Lack customizability for topologies, forwarding behaviors, metrics (link delay, link loss, ...)
- 😞 Costly to scale and upgrade

***Scalability & Flexibility Concerns***

# To conduct network experiments



	Simulator	Emulator	Test-Bed
Fidelity	×	×	✓
Flexibility	✓	✓	×
Scale	✓	×	×

# Our Approach

**TurboNet: emulating various networks with one programmable switch.**

## *Why programmable switch?*

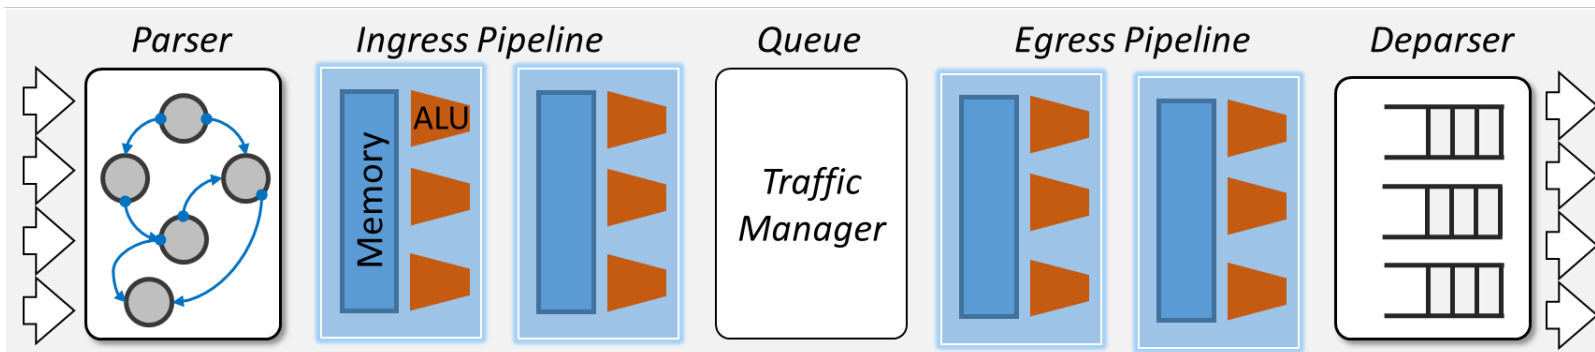
- ❑ Re-configurable packet processing
- ❑ High bandwidth (Tbps)

**Enable agile customization of various networks with high bandwidth**

## *Our Goal:*

**Faithfully emulate both network data plane and control plane.**

- ❑ Data Plane: topology, metrics (delay, loss)
- ❑ Control Plane: routing (static/dynamic)



[RMT@SIGCOMM'13]

# Talk Outline

- ~~Motivation~~
- **1. TurboNet Overview: architecture and workflow**
- 2. Data Plane Emulation
- 3. Control Plane Emulation
- 4. Some Evaluation Results
- 5. Conclusion

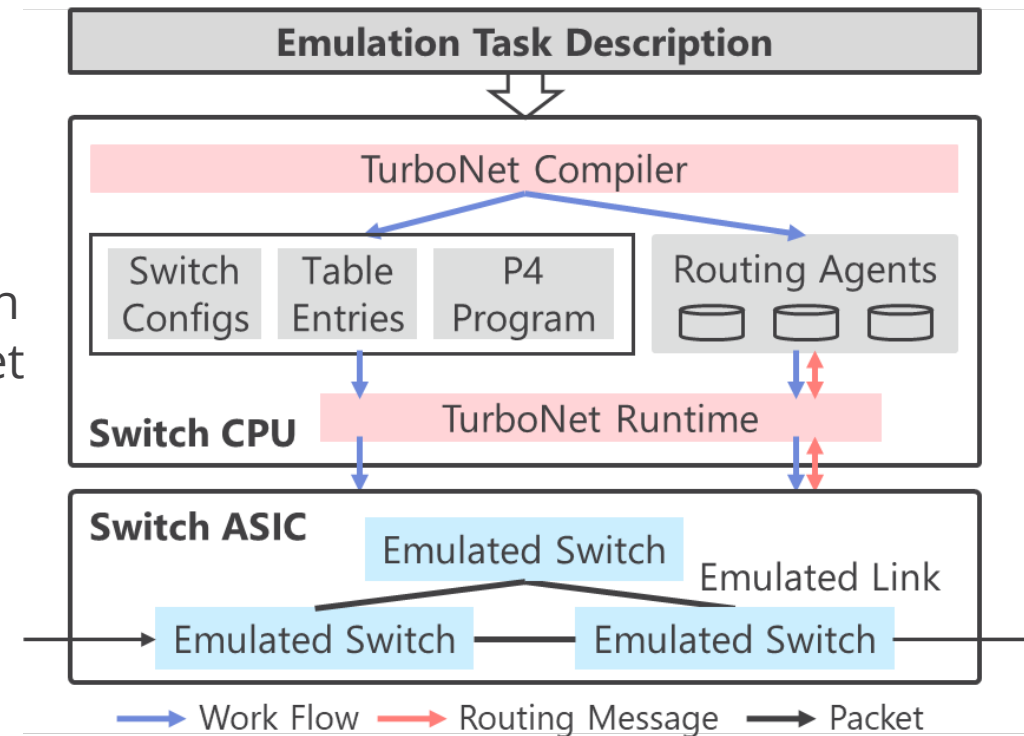


# TurboNet Overview




A programmable switch  
deployed with TurboNet

## TurboNet Architecture



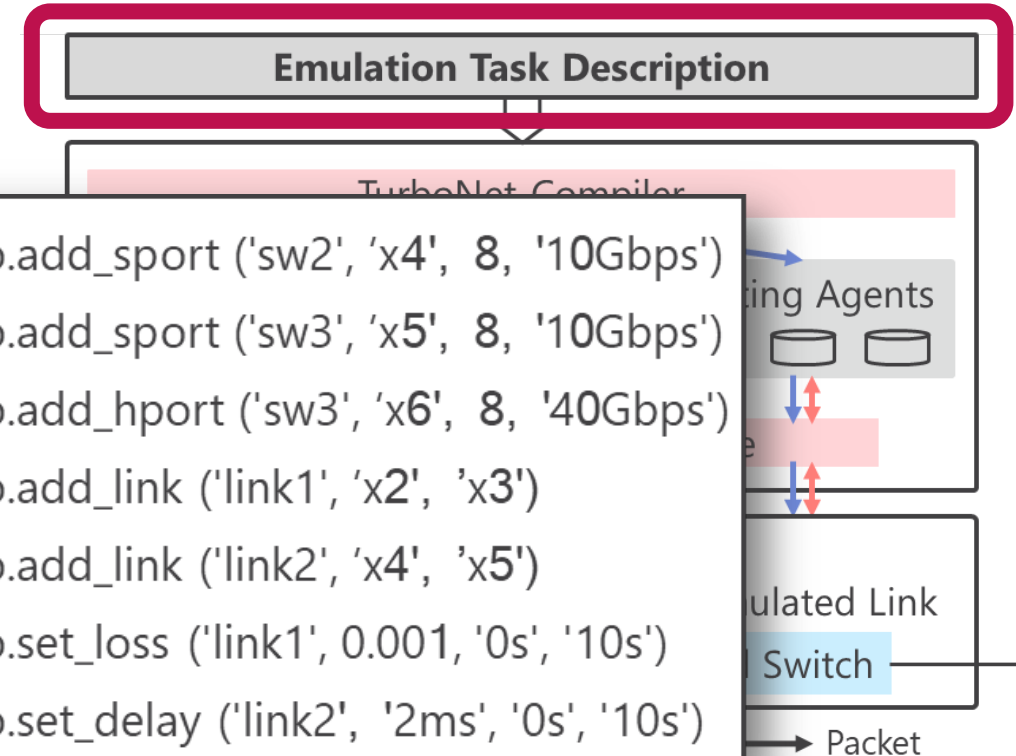
## TurboNet Overview

### 1. Emulation Task Description




```
1 topo = Topo ()
2 topo.add_switch ('sw1')
3 topo.add_switch ('sw2')
4 topo.add_switch ('sw3')
5 topo.add_hport ('sw1', 'x1', 8, '40Gbps')
6 topo.add_sport ('sw1', 'x2', 8, '10Gbps')
7 topo.add_sport ('sw2', 'x3', 8, '10Gbps')
8 topo.add_sport ('sw2', 'x4', 8, '10Gbps')
9 topo.add_sport ('sw3', 'x5', 8, '10Gbps')
10 topo.add_hport ('sw3', 'x6', 8, '40Gbps')
11 topo.add_link ('link1', 'x2', 'x3')
12 topo.add_link ('link2', 'x4', 'x5')
13 topo.set_loss ('link1', 0.001, '0s', '10s')
14 topo.set_delay ('link2', '2ms', '0s', '10s')
```

### TurboNet Architecture



# TurboNet Workflow



# 1. Emulation Task Description



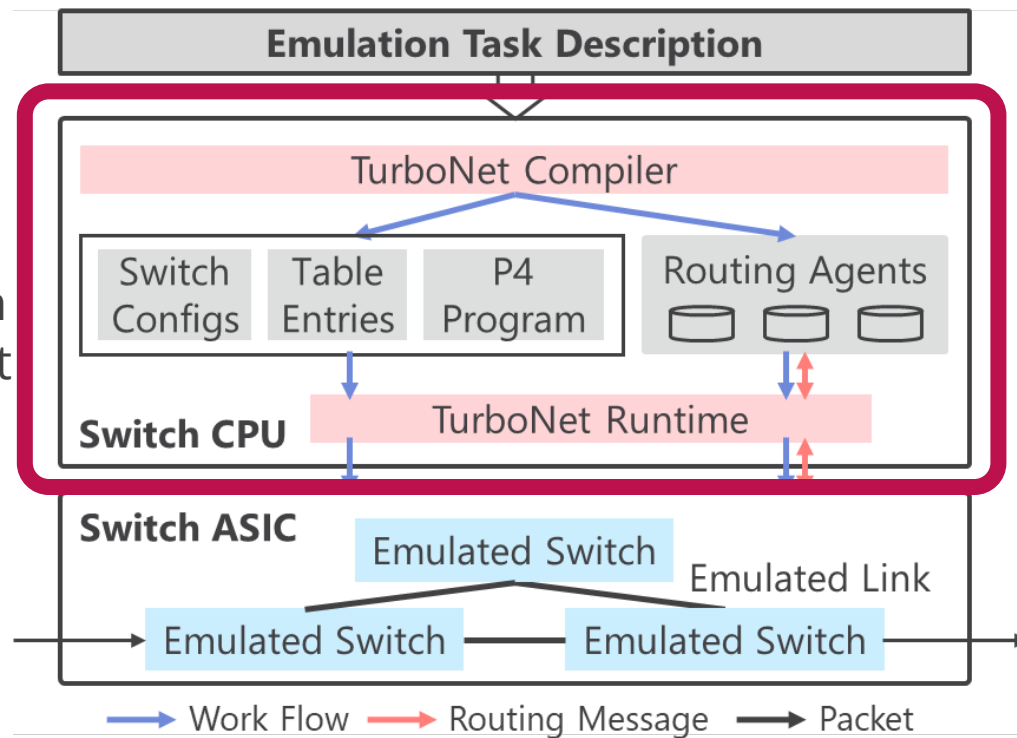
## 2. Configuration Generation




## A programmable switch deployed with TurboNet

# TurboNet Overview

# TurboNet Architecture




# TurboNet Workflow




## 1. Emulation Task Description



## 2. Configuration Generation

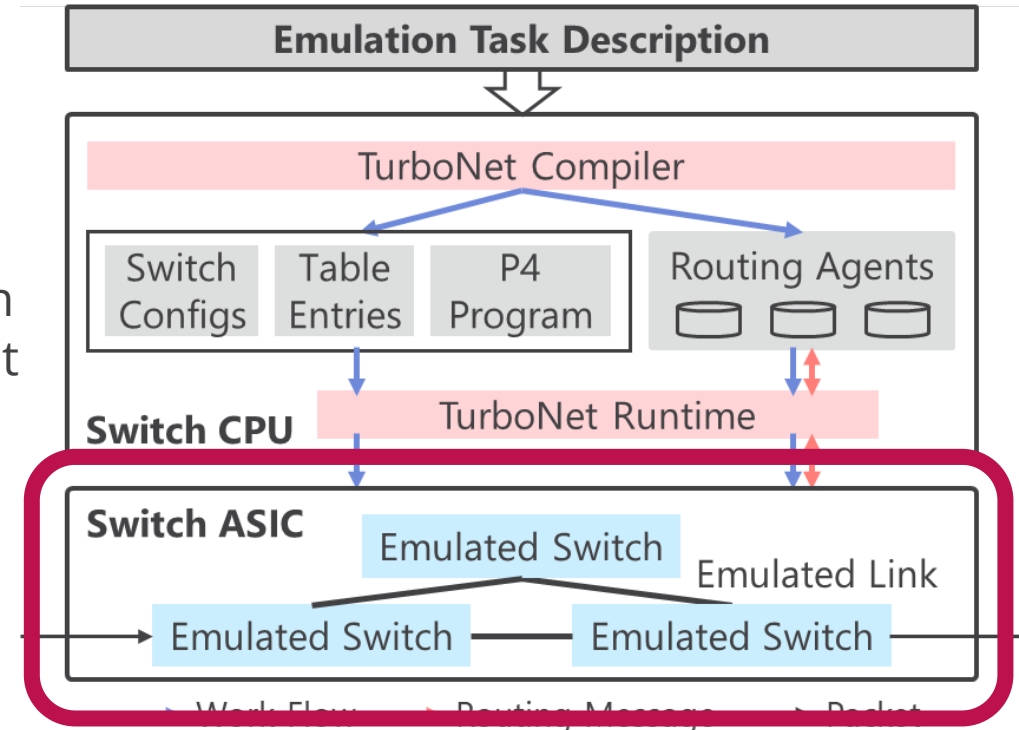


A programmable switch deployed with TurboNet




# TurboNet Overview


## TurboNet Architecture



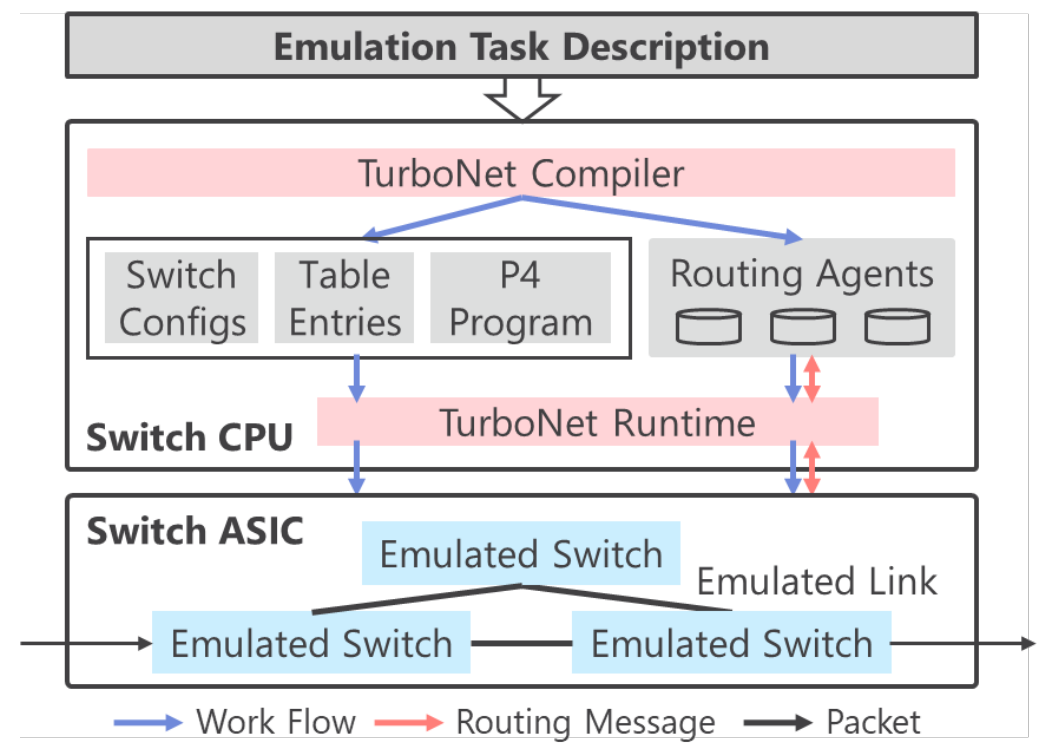
# TurboNet Overview



How?



## TurboNet Architecture



# Talk Outline

- Motivation
- 1. TurboNet Overview
- **2. Data Plane Emulation**
- 3. Control Plane Emulation
- 4. Some Evaluation Results
- 5. Conclusion

# Data Plane Emulation

## Forwarding packets in a network with specified performance

- Topology Emulation: flexibly emulate various topologies in a programmable switch
- Metric Emulation: make the emulated network behave like real networks
  - e.g., link loss, link delay, ...

# Data Plane Emulation

## Forwarding packets in a network with specified performance

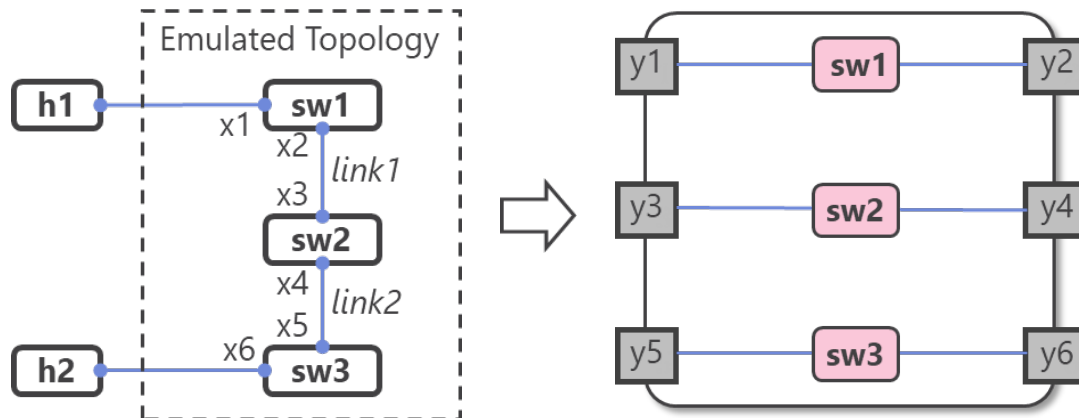
- Topology Emulation: flexibly emulate various topologies in a programmable switch
- Metric Emulation: make the emulated network behave like real networks
  - e.g., link loss, link delay, ...



# Topology Emulation

**Key Idea: Split one programmable switch into multiple emulated**

## □ Port Mapper



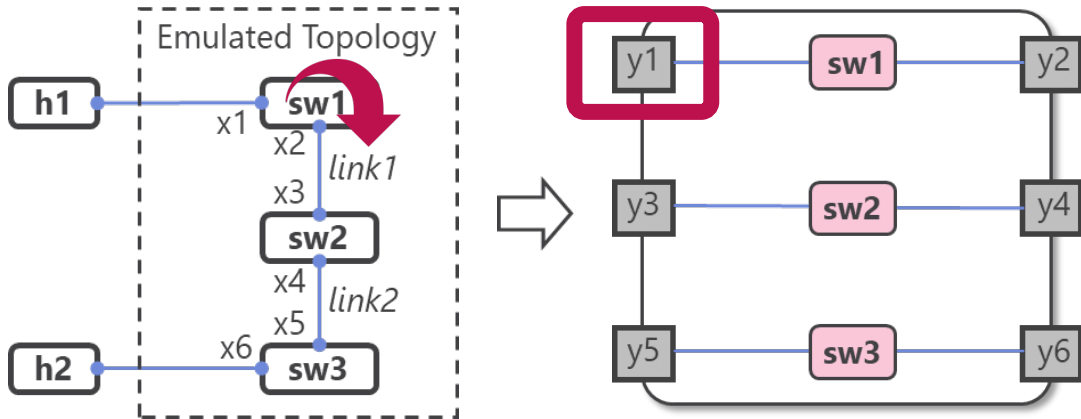
*Map each port in the input topology to an individual port on the programmable switch ( x1~x6 -> y1~y6)*

**Switch  
Emulation**


# Topology Emulation

## Key Idea: Split one programmable switch into multiple emulated

## □ Port Mapper



**Map each port in the input topology to an individual port on the programmable switch ( x1~x6 -> y1~y6)**

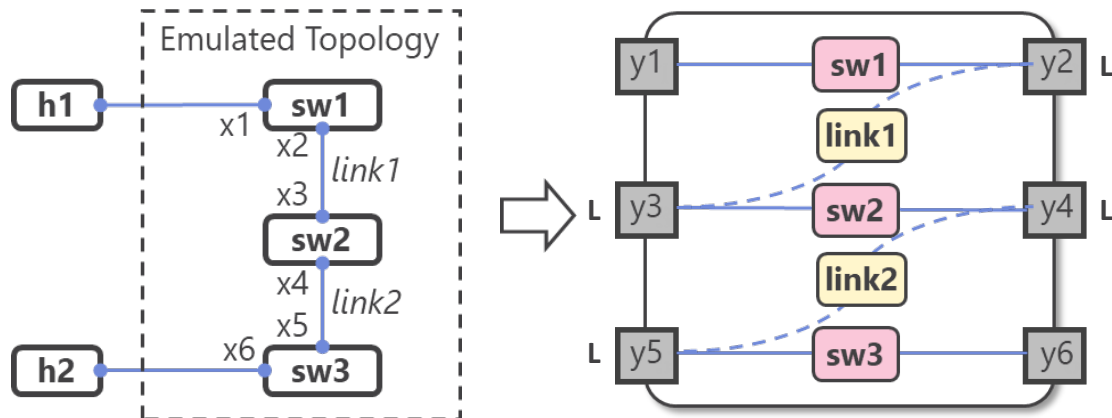


# Switch Emulation


# Topology Emulation

**Key Idea: Split one programmable switch into multiple emulated**

## □ Port Mapper



*Map each port in the input topology to an individual port on the programmable switch (x1~x6 -> y1~y6)*



## Switch Emulation

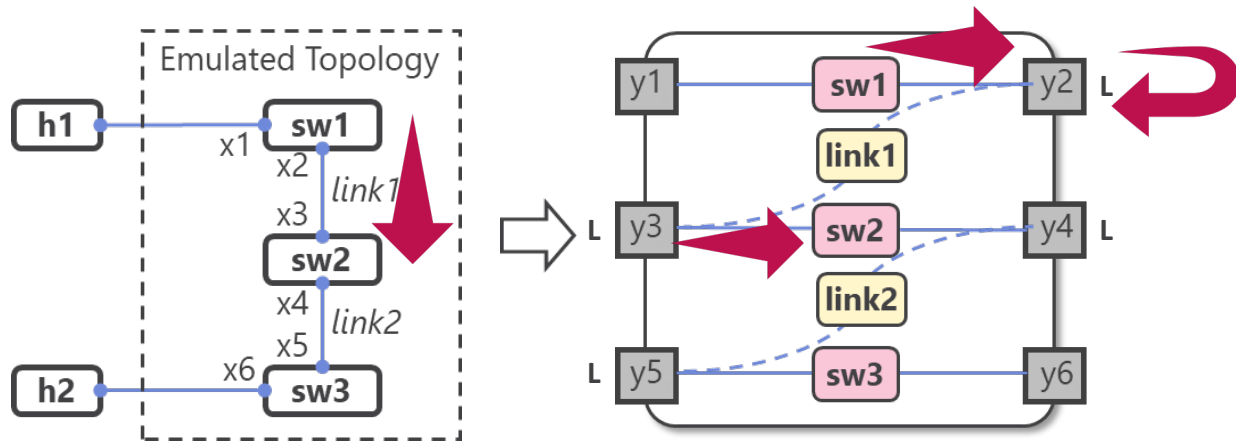
### Link Emulation

- Cables
- Loopback

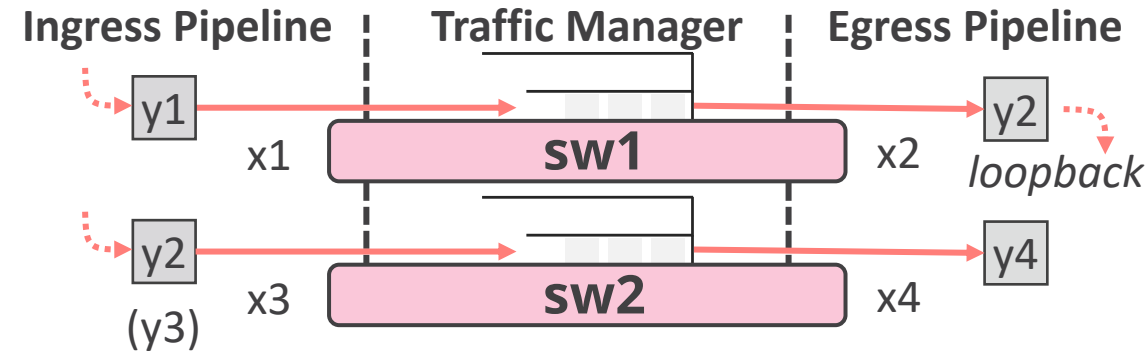
# Topology Emulation

**Key Idea: Split one programmable switch into multiple emulated**

## □ Port Mapper



*Map each port in the input topology to an individual port on the programmable switch (x1~x6 -> y1~y6)*



## Switch Emulation

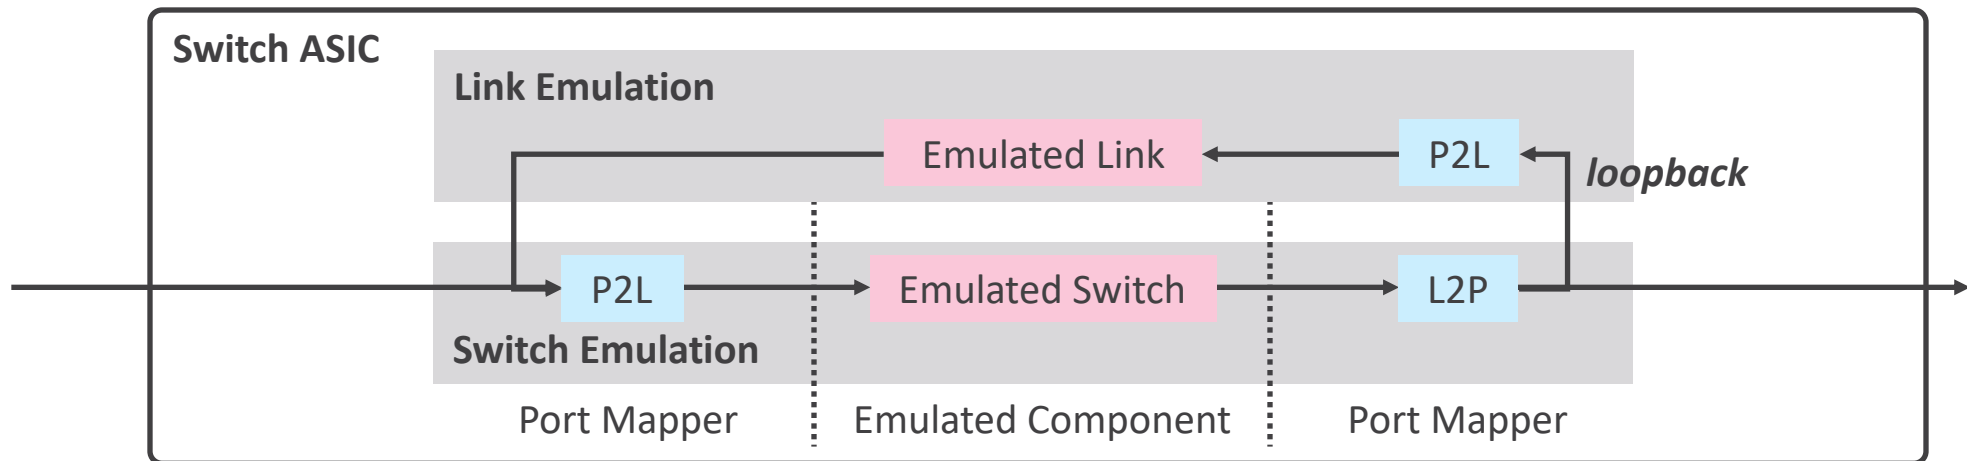
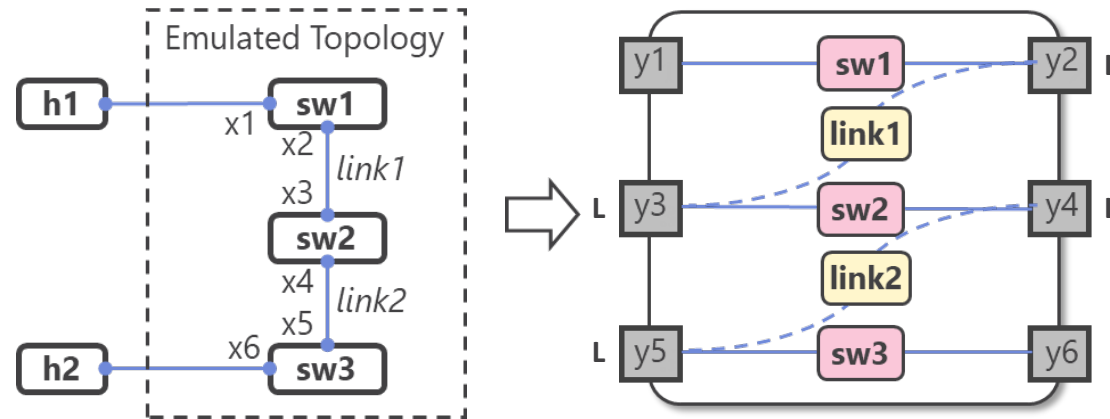
### Link Emulation

- Cables
- Loopback

# Topology Emulation

**Key Idea: Split one programmable switch into multiple emulated**

## □ Port Mapper



# Topology Emulation

**Key Idea: Split one programmable switch into multiple emulated**

□ Port Mapper *At most 64\*4 ports?*

Barefoot Tofino switch

- 64 100Gbps port groups
- A port group:
  - 1x100G, 2x40G, 4x25G, 4x10G ports

# Topology Emulation

**Key Idea: Split one programmable switch into multiple emulated**

❑ Port Mapper *At most 64\*4 ports?*

**Challenge: How to emulate larger networks?**

❑ Queue Mapper! *At most 64\*32 queues!*

Barefoot Tofino switch

- 64 100Gbps port groups
- A port group:
  - 1x100G, 2x40G, 4x25G, 4x10G ports
  - 32 queues

# Topology Emulation

**Key Idea: Split one programmable switch into multiple emulated**

□ Port Mapper *At most 64\*4 ports?*

**Challenge: How to emulate larger networks?**

□ Queue Mapper! *At most 64\*32 queues!*

**Problem: How to perform port/queue mapping?**

□ *0-1 ILP problem*



# Topology Emulation

**Key Idea: Split one programmable switch into multiple emulated**

□ Port Mapper *At most 64\*4 ports?*

**Challenge: How to emulate larger networks?**

□ Queue Mapper! *At most 64\*32 queues?*

**Problem: How to perform**

□ **0-1 ILP problem**

Input {  
Ports  
Bandwidth  
Queue number

*(See more in our paper)*

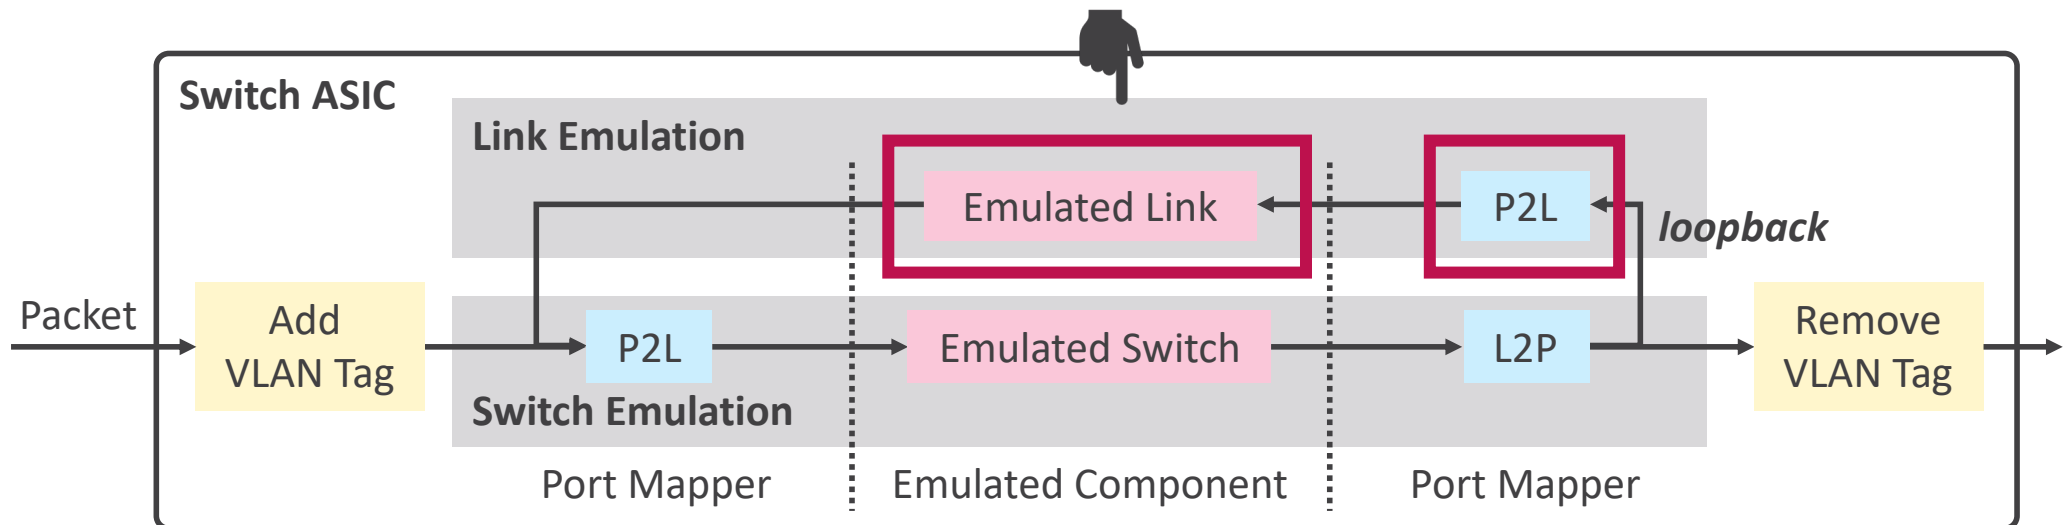
PM Problem	$x_i^{1,2,3,4}, x_{ik}^{1,2,3,4} = PM(S, b_k, q_k, P)$
Objective	$\min: \sum_{i \in P} (x_i^1 + x_i^2 + x_i^3 + x_i^4)$
	$3 + x_i^4 \leq 1, \forall i \in P$
	$3 \sum x_{ik}^4 \leq 4 \cdot x_i^4, \forall i \in P$
QM Problem	$y_i, y_{ik}, x_i^{1,2,3,4}, x_{ik}^{1,2,3,4} = QM(S_1, S_2, b_k, q_k, P)$
Objective	$\min: \sum_{i \in P} x_i^1 + x_i^2 + x_i^3 + x_i^4 + y_i$
Constraints	$C_1: x_i^1 + x_i^2 + x_i^3 + x_i^4 + y_i \leq 1, \forall i \in P$ $C_2: x_i^{1,2,3,4}, x_{ik}^{1,2,3,4} = PM(S_2, b_k, q_k, P)$ $C_3: \sum_{k \in S_1} b_k y_{ik} \leq 100 y_i, \forall i \in P$ $C_4: \sum_{i \in P} y_{ik} = 1, \forall k \in S_1$ $C_5: \sum_{k \in S_1} q_k y_{ik} \leq 32 y_i, \forall i \in P$

# Data Plane Emulation

- Topology Emulation
- Metric Emulation: make the emulated network behave like real networks
  - Link loss
  - Link delay
- Background traffic

# Link Loss Emulation

## *Link Metric Emulation*



# Link Loss Emulation

Generate a random number  $v$

Compare  $v$  with threshold  $\theta$

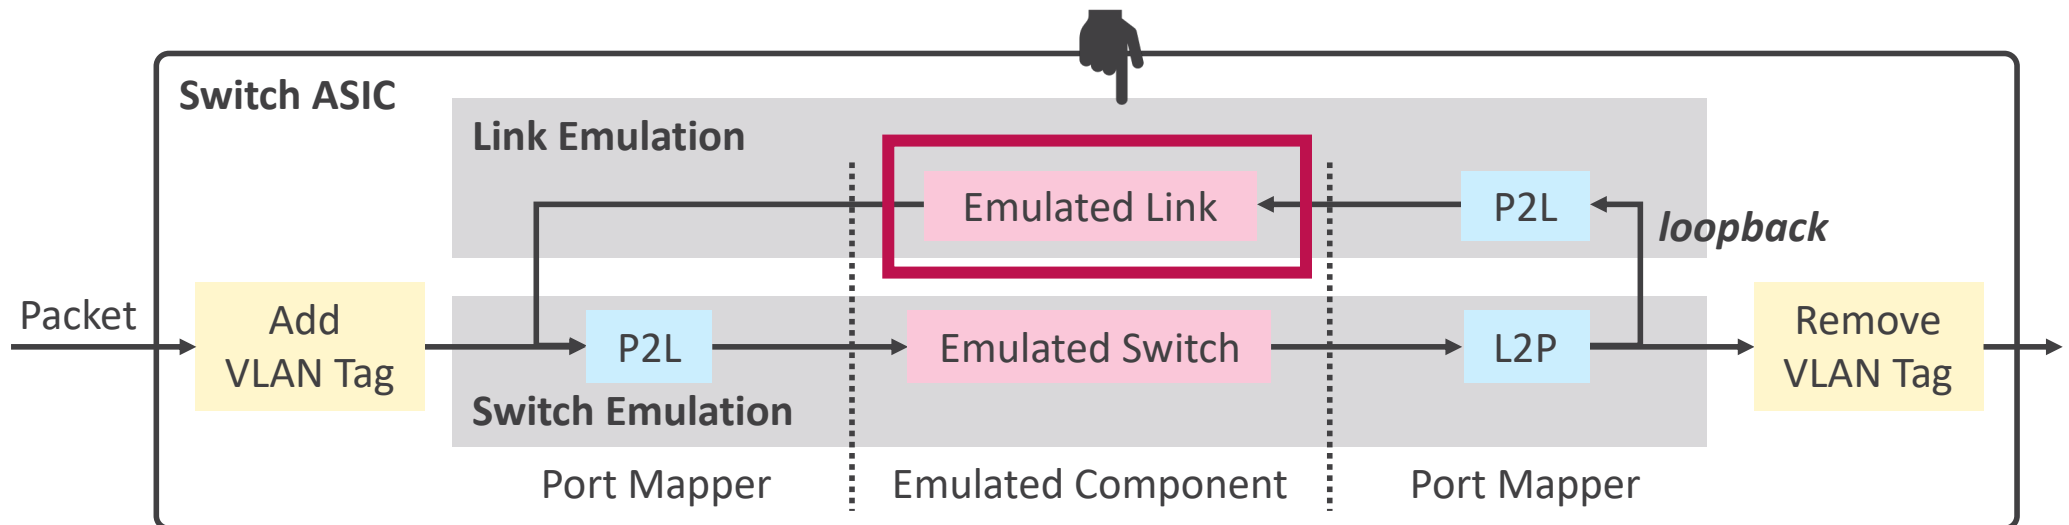
Drop if  $v > \theta$

P4 primitive: `modify_field_rng_uniform`


$\theta$  is decided by given loss rate  $r$ :  $P(v > \theta) = r$

Drop packets according to a probability

## Link Loss Emulation



# Link Delay Emulation

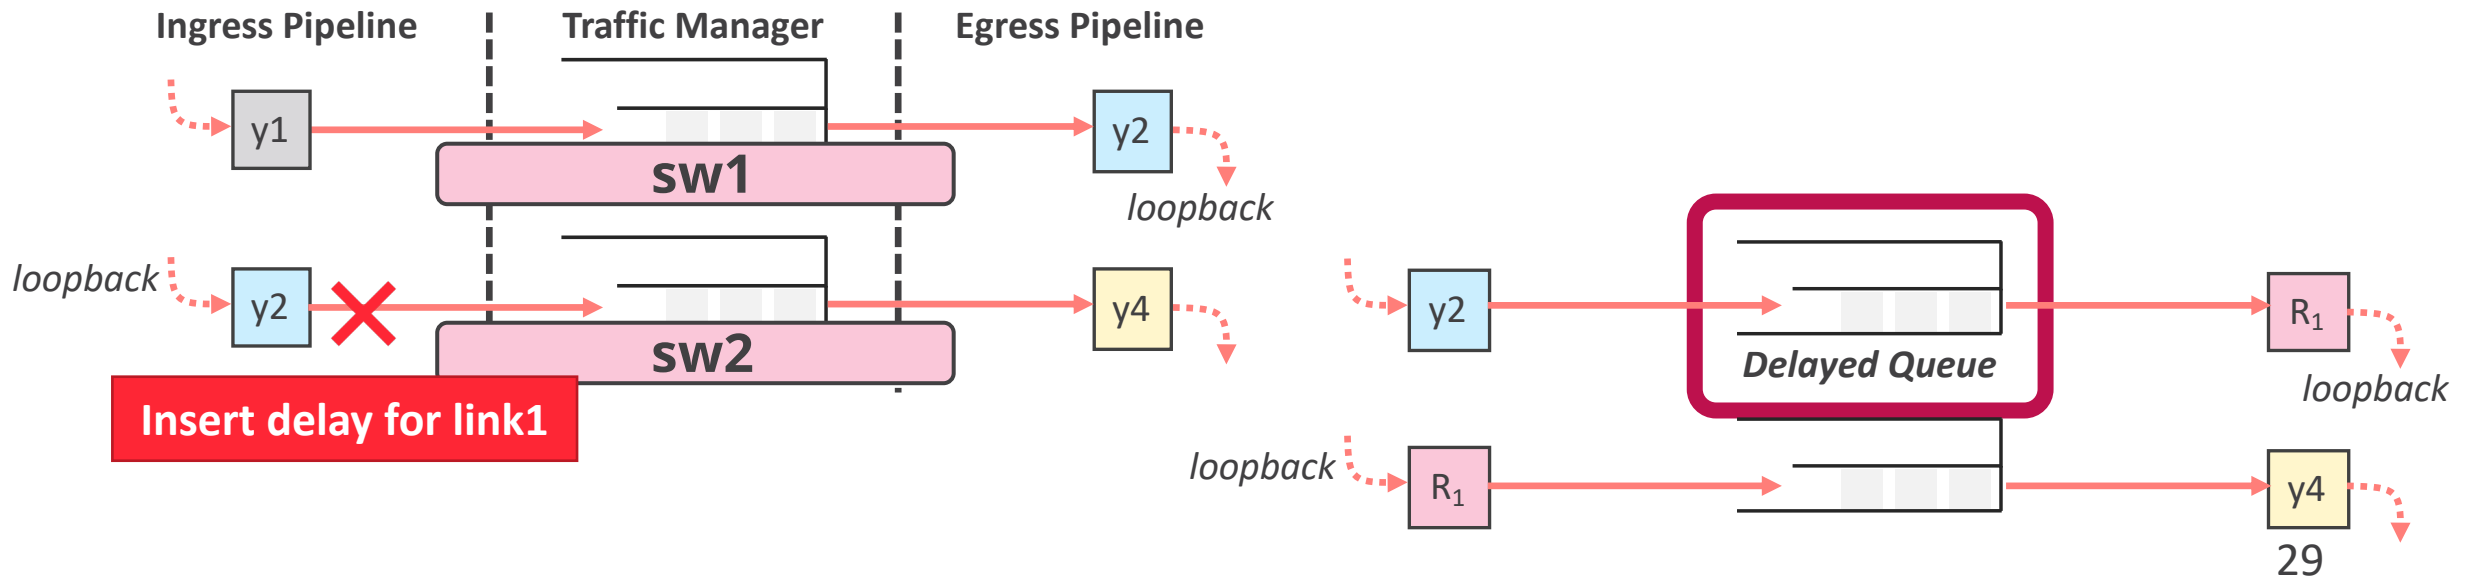


**Challenge: How to keep packets in programmable switch?**

**Observation: queuing time is in proportion to queue depth.**

**To achieve target queuing time, maintain queue depth!**

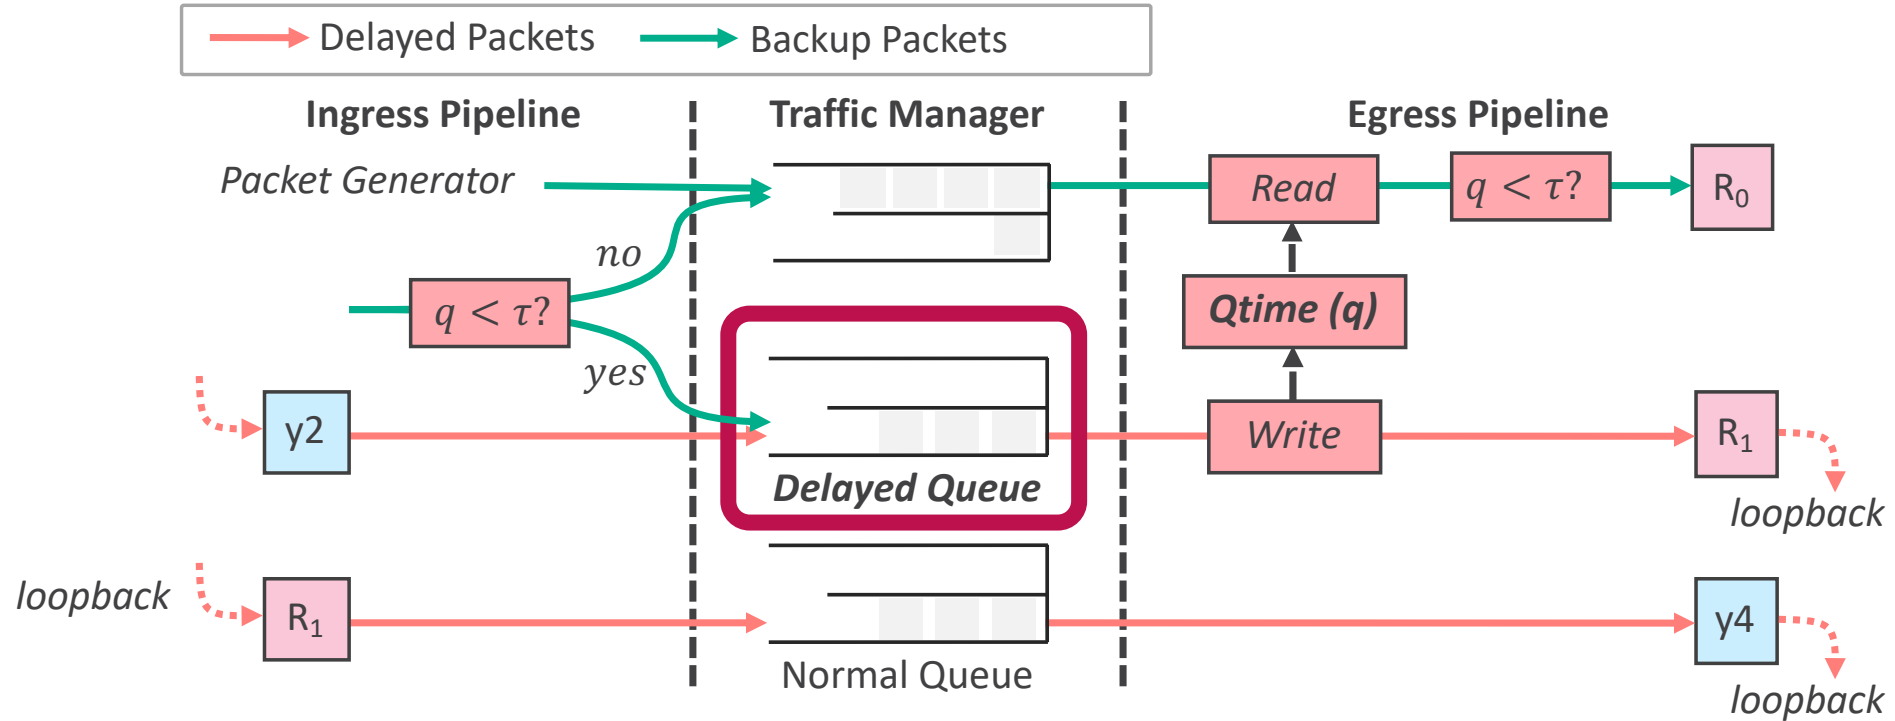
- ❑ Implement **delayed queues** with desired queuing time
- ❑ Send packets to delayed queues to emulate target delay



# Link Delay Emulation

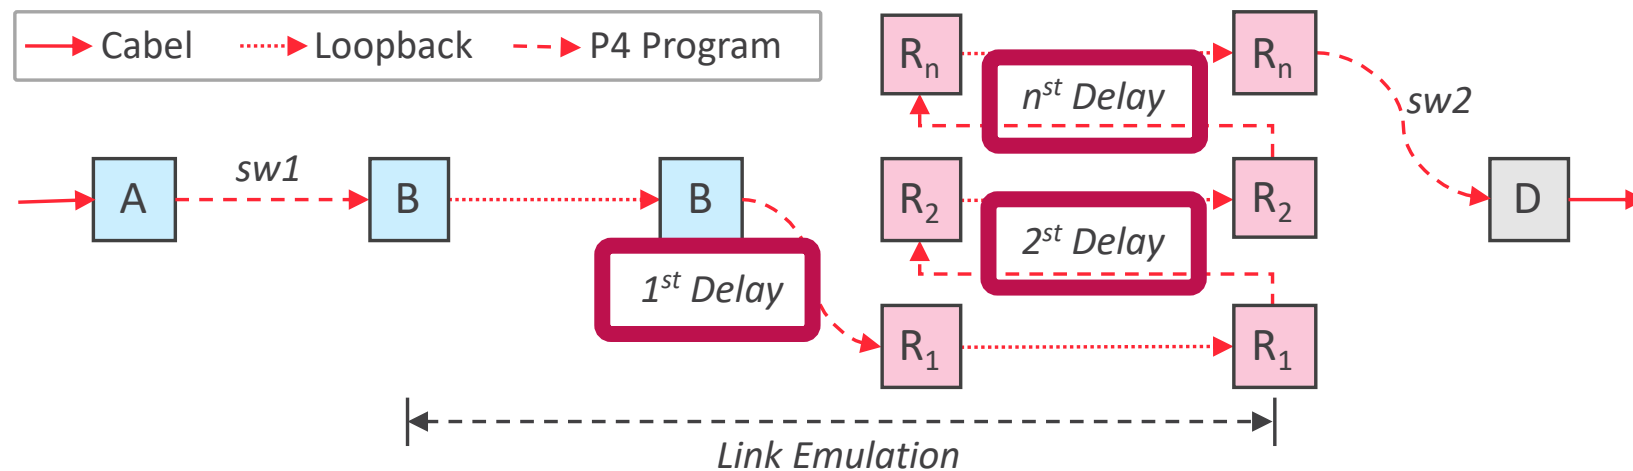
## 1. Implement delayed queues with desired queuing time

- dynamically inject packets



# Link Delay Emulation

1. Implement delayed queues with desired queuing time
  - ▣ dynamically inject backup packets
2. Send packets to delayed queues to emulate target delay



***More details (resource, accuracy) in our paper***

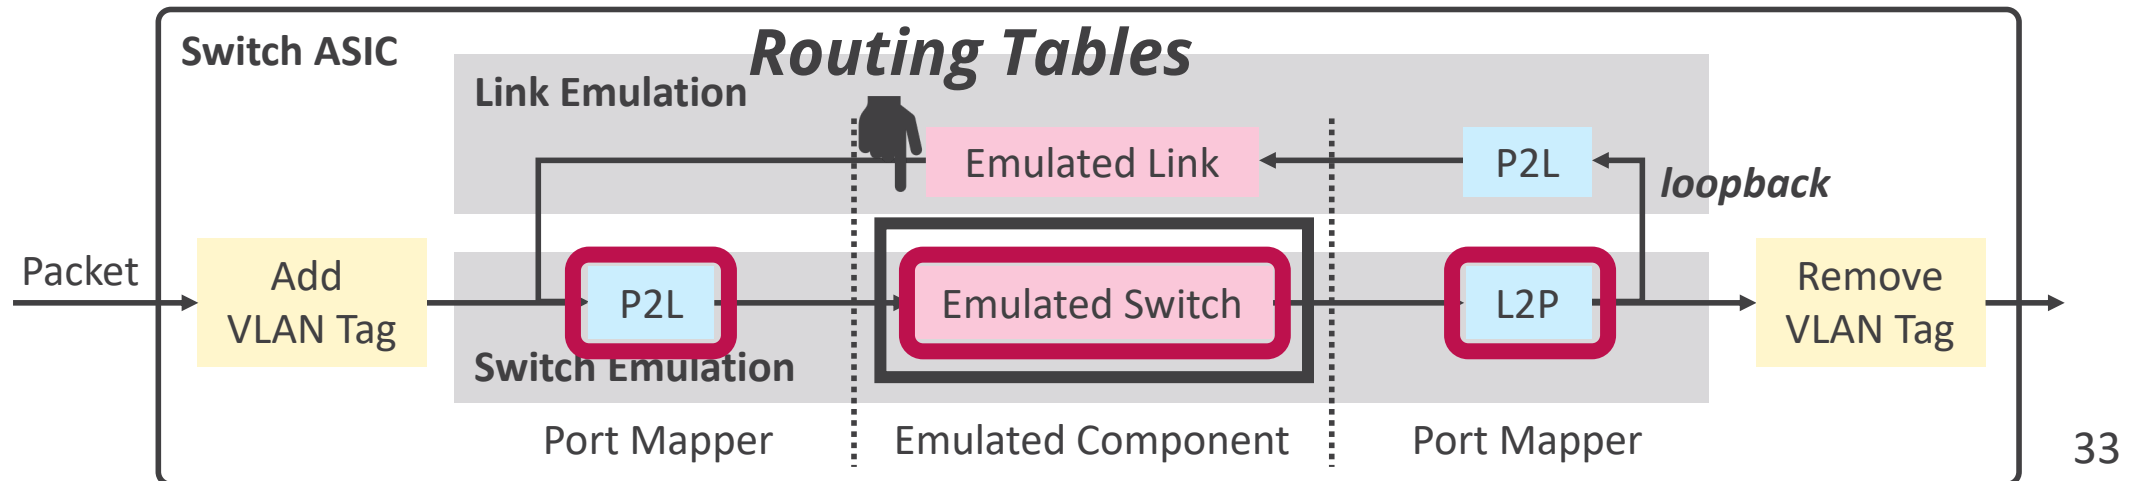
# Talk Outline

- Motivation
- 1. TurboNet Overview
- 2. Data Plane Emulation
- **3. Control Plane Emulation**
- 4. Some Evaluation Results
- 5. Conclusion



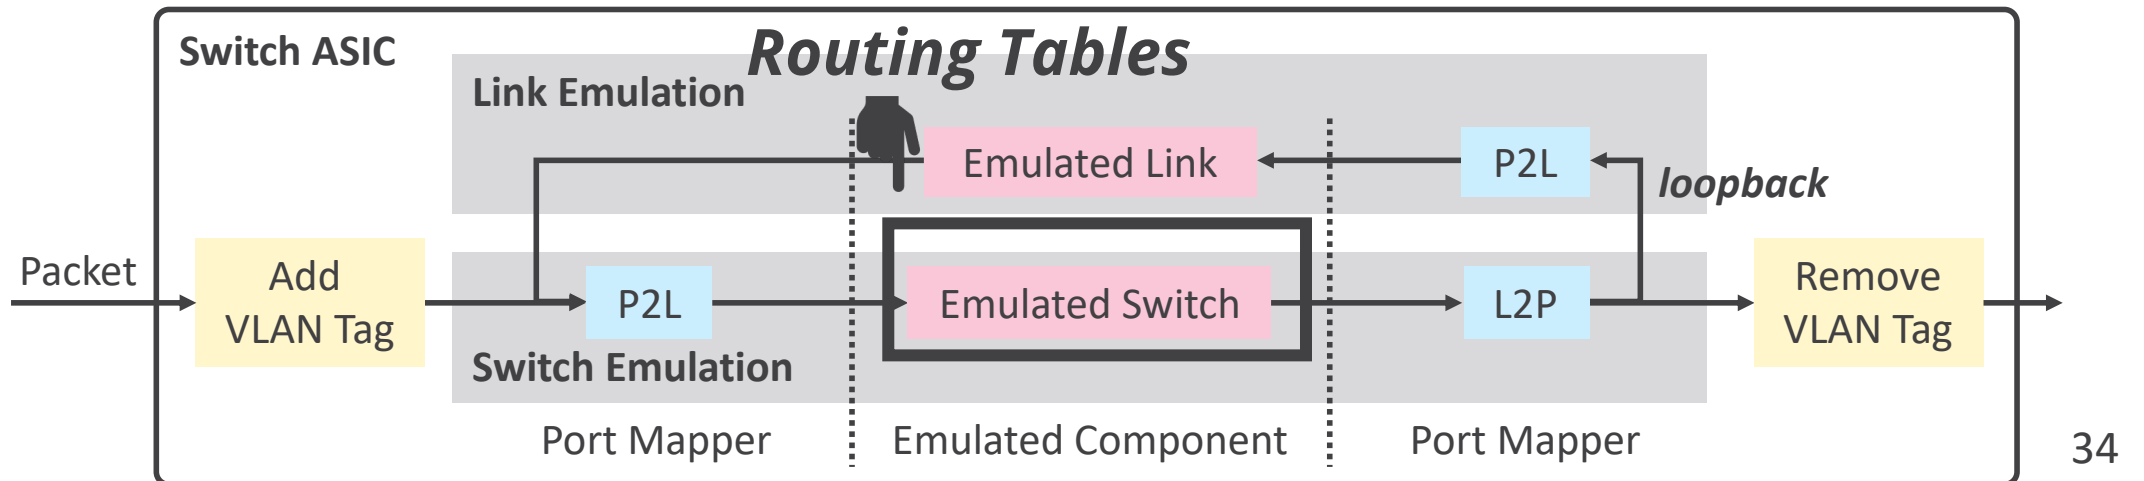
# Control plane emulation

- Static Routing
- Dynamic routing



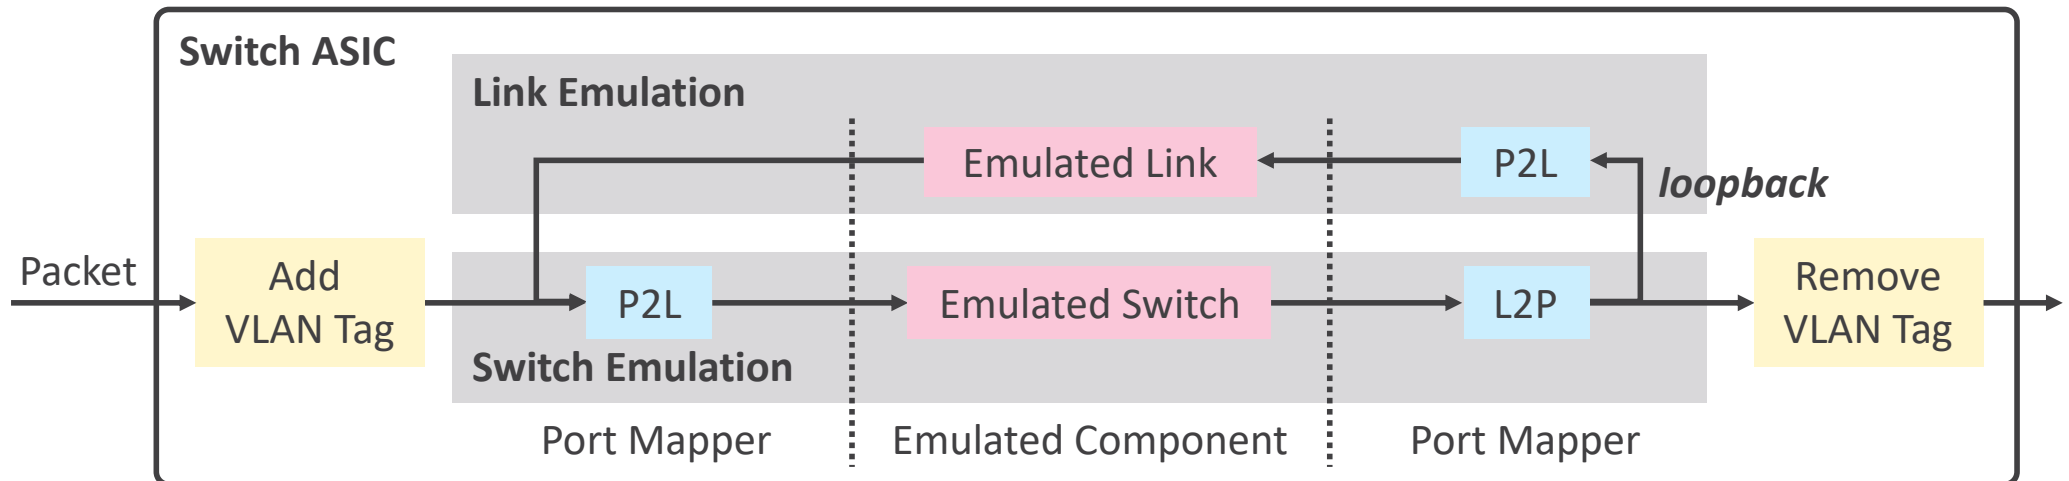
# Control plane emulation

- Static Routing
- Dynamic routing




# Control plane emulation

- Static Routing
- **Dynamic routing**
  - Distributed routing
  - Centralized routing



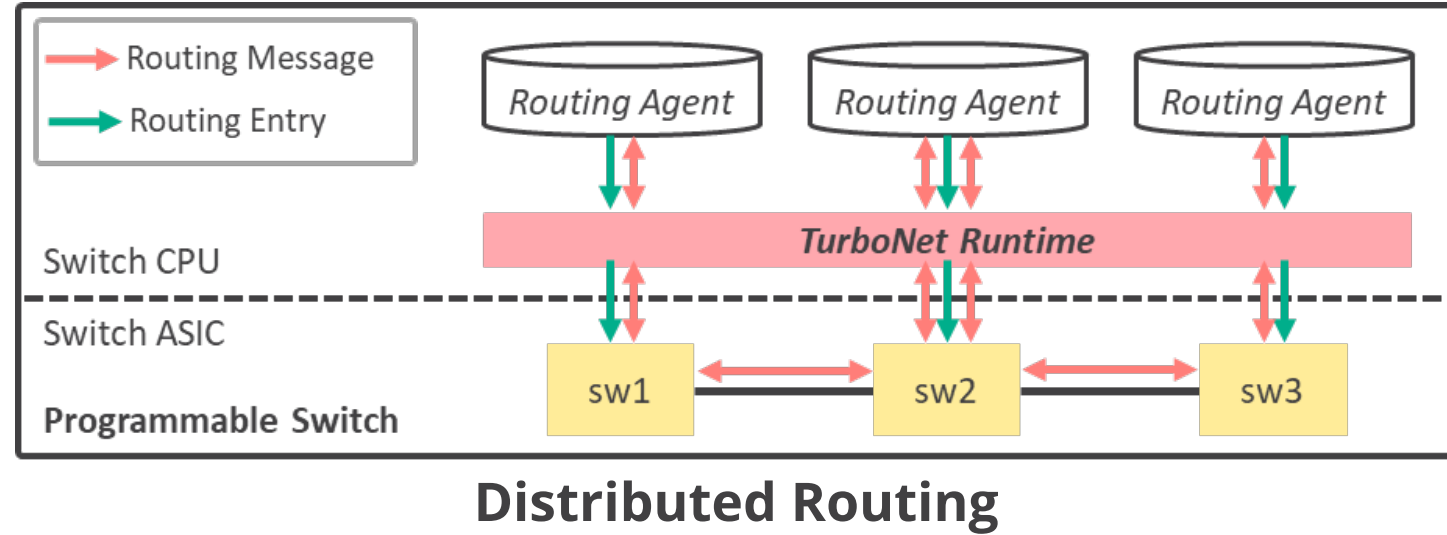
# Control plane emulation

- Static Routing
- Dynamic routing
  - Distributed routing
  - Centralized routing




# Control plane emulation

- Static Routing
- Dynamic routing
  - Distributed routing
  - Centralized routing

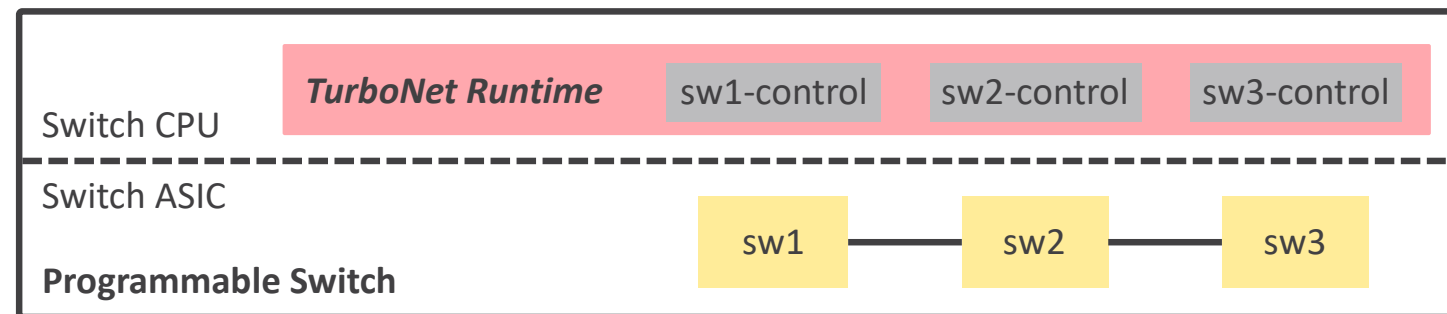
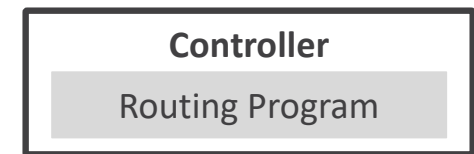


# Control plane emulation

- Static Routing
- Dynamic routing
  - Distributed routing
  - Centralized routing



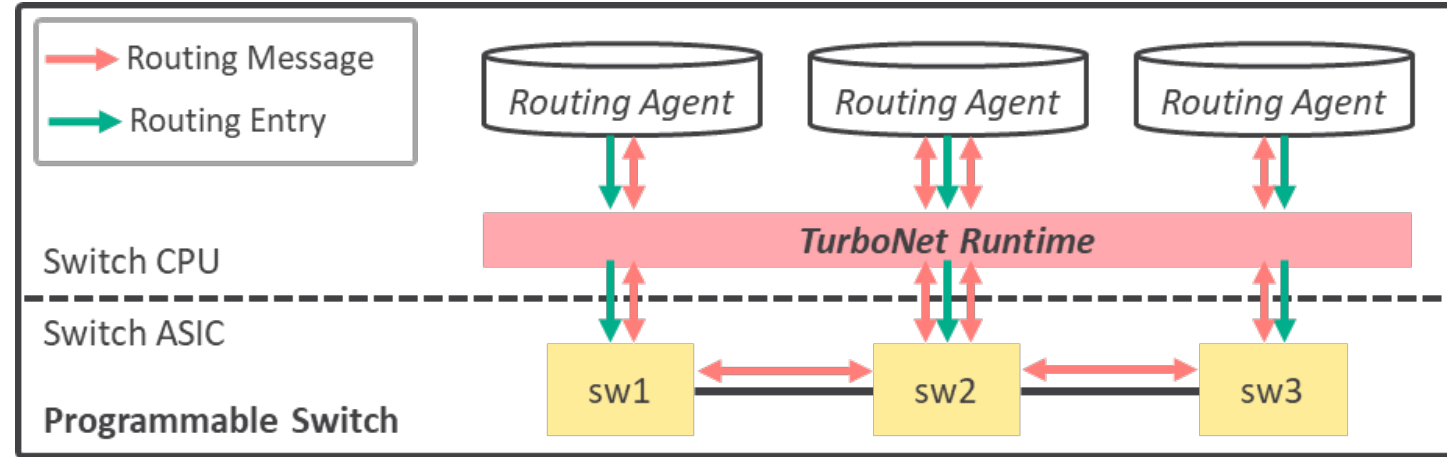
**Distributed Routing**



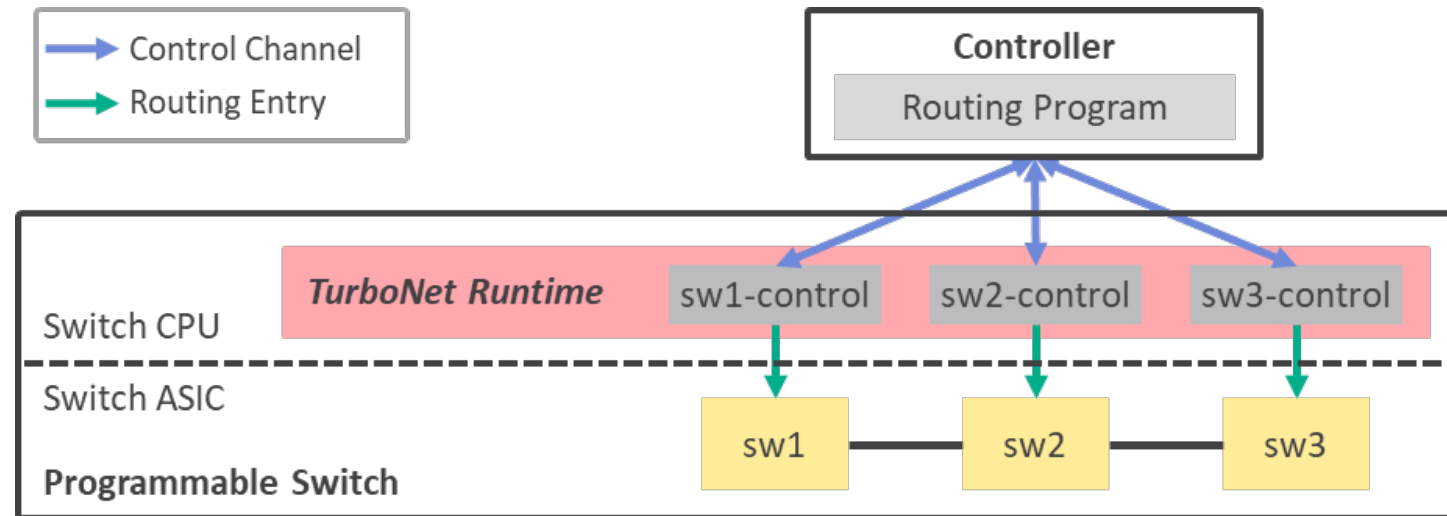
**Centralized Routing**

# Control plane emulation

- Static Routing
- Dynamic routing
  - Distributed routing
  - Centralized routing



**Distributed Routing**



**Centralized Routing**

# Talk Outline

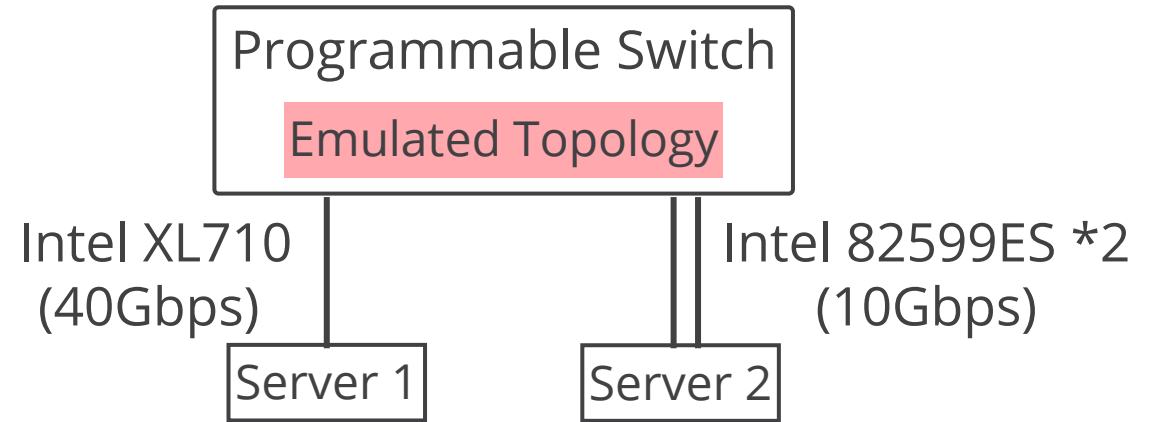
- Motivation
- 1. TurboNet Overview
- 2. Data Plane Emulation
  - Topology Emulation
  - Metric Emulation: Background Traffic, Link Loss, Link Delay
- 3. Control Plane Emulation
- **4. Some Evaluation Results**
- 5. Conclusion



# Evaluation Setup

## Test-bed

- One Barefoot Tofino Switch
  - Intel Pentium 4-core 1.60GHz CPU
  - 8GiB memory
  - 2 pipelines (3.2Tbps)
- Two servers
  - 12-core Intel Xeon E5-2620  
2.40GHz CPU



# Topology Emulation Capability

Topology		TurboNet			
		2-Pipe (32*100G) Programmable Switch		4-Pipe (64*100G) Programmable Switch	
		PM	QM	PM	QM
Fat-Tree	$k = 4$	$Link \leq 25G$	$Link \leq 25G$	$Link \leq 40G$	$Link \leq 40G$
	$k = 6$	$\times$	$Link \leq 8.3G$	$\times$	$Link \leq 20G$
	$k = 8$	$\times$	$\times$	$\times$	$Link \leq 6.2G$
VL2	$Link = 1G, 10G$	$D_A D_1 \leq 16$	$D_A D_1 \leq 20$	$D_A D_1 \leq 36$	$D_A D_1 \leq 44$
	$Link = 10G, 40G$	$D_A D_1 \leq 20$	$D_A D_1 \leq 20$	$D_A D_1 \leq 40$	$D_A D_1 \leq 40$
261 Internet Topologies		199 (76.2%)	259 (99.2%)	248 (95.0%)	260 (99.6%)

## PM VS. QM

- Fat-tree:  $k=4$  for PM,  $k=8$  for QM

Note: A  $k$ -ary fat-tree has  $k^3/4$  host-ports and  $k^3$  switch-ports<sup>42</sup>

# Topology Emulation Capability


Topology		TurboNet			
		2-Pipe (32*100G) Programmable Switch		4-Pipe (64*100G) Programmable Switch	
		PM	QM	PM	QM
Fat-Tree	$k = 4$	$Link \leq 25G$	$Link \leq 25G$	$Link \leq 40G$	$Link \leq 40G$
	$k = 6$	$\times$	$Link \leq 8.3G$	$\times$	$Link \leq 20G$
	$k = 8$	$\times$	$\times$	$\times$	$Link \leq 6.2G$
VL2	$Link = 1G, 10G$	$D_A D_1 \leq 16$	$D_A D_1 \leq 20$	$D_A D_1 \leq 36$	$D_A D_1 \leq 44$
	$Link = 10G, 40G$	$D_A D_1 \leq 20$	$D_A D_1 \leq 20$	$D_A D_1 \leq 40$	$D_A D_1 \leq 40$
261 Internet Topologies		199 (76.2%)	259 (99.2%)	248 (95.0%)	260 (99.6%)

## PM VS. QM

- Fat-tree:  $k=4$  for PM,  $k=8$  for QM
- Internet topology: 60 more for QM

Note: A  $k$ -ary fat-tree has  $k^3/4$  host-ports and  $k^3$  switch-ports<sup>43</sup>


# Link Loss Accuracy



Link loss accuracy  
(32-bit random number generation)


Expected loss rate changes from 0.1 to  $10^{-8}$   
Relative error  $< 1\%$

# Link Delay Accuracy



One-pass link delay  
(A 10Gbps delayed queue)

Expected delay ranges from 0.2ms to 1ms  
Standard deviation < 9 $\mu$ s





Multi-pass link delay  
(A 1ms delayed queue)

Expected delay ranges from 2ms to 10ms  
Relative standard deviation < 0.9%

# Control Plane Emulating Capability

Intel Pentium 4-core 1.60GHz CPU, 8GiB memory  
All connected routers are BGP peers



TurboNet can easily support almost 200 BGP routing agents with 25% peak CPU usage.

# Talk Outline

- Motivation
- 1. TurboNet Overview
- 2. Data Plane Emulation
  - Topology Emulation
  - Metric Emulation: Background Traffic, Link Loss, Link Delay
- 3. Control Plane Emulation
- **4. Some Evaluation Results**
- **5. Conclusion**

# Conclusion

- **TurboNet is a network experiment platform that leverages one programmable switch to faithfully mimic networks.**
  - Data plane
  - Control plane



# Thanks!

